

Another Linux Kernel Bug Surfaces, Allowing Root Access



Author:
Tara Seals

September 28, 2018
/ 2:11 pm



October 7, 2019

Kernel privilege escalation bug actively exploited in Android devices

Bradley Barth

[Follow @bbb1216bbb](#)

Low Level Software Security

Computer Security and Privacy (CS642)

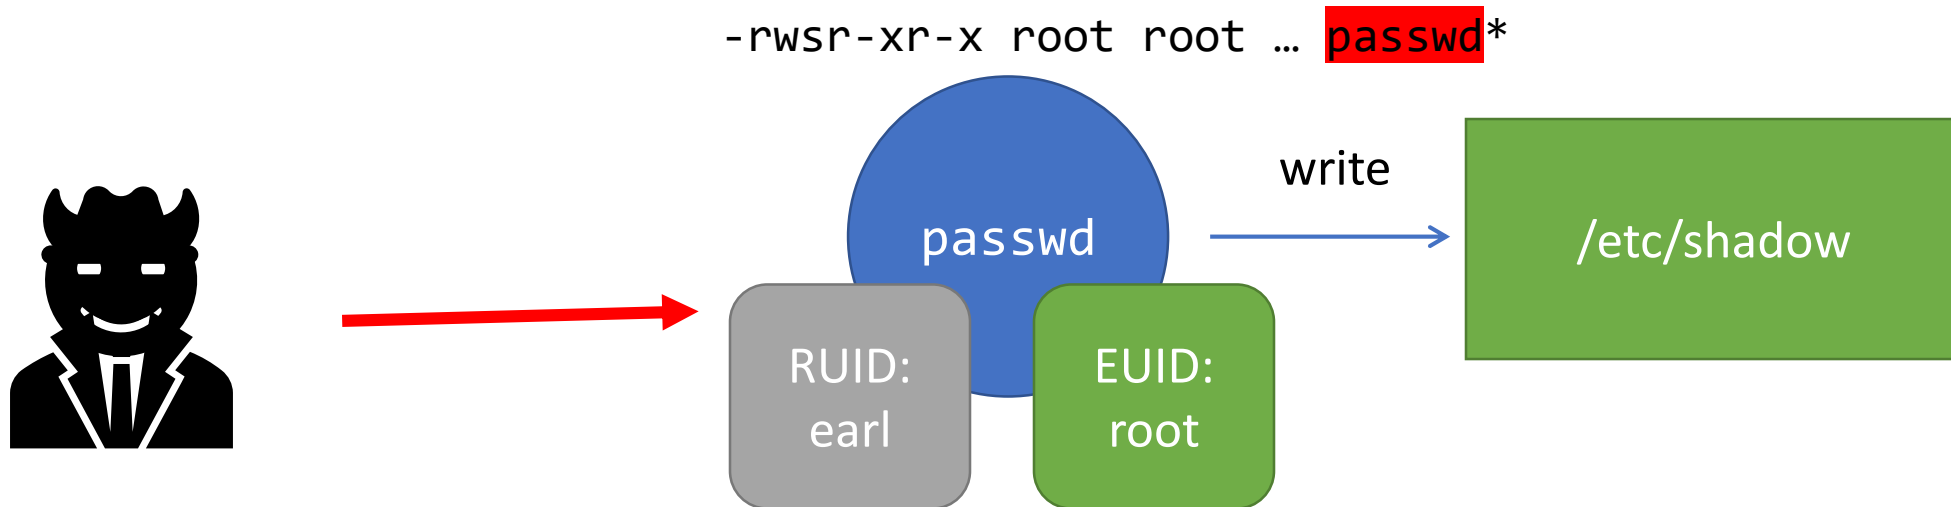
Earlence Fernandes

earlence@cs.wisc.edu

* Slides borrowed from Chatterjee, Davidson, Ristenpart


Processes are the front line of system security

- Control a process and you get the privileges of its UID
- So how do you control a process?
 - Send specially formed input to process



Roadmap

- Today
 - Enough x86 to understand (some) process vulnerabilities
 - ISA
 - Process memory layout, call stack
 - Buffer overflow attack
 - How such attacks occur



Why do we need to look at assembly?

WYSINWYX: What You See Is Not What You eXecute

G. Balakrishnan¹, T. Reps^{1,2}, D. Melski², and T. Teitelbaum²

¹ Comp. Sci. Dept., University of Wisconsin; {bgogul,reps}@cs.wisc.edu

² GrammaTech, Inc.; {melski,tt}@grammatech.com

We understand code in this form

```
int foo() {  
    int a = 0;  
    return a + 7;  
}
```



Compiler

Vulnerabilities exploited in this form

```
pushl %ebp  
movl  %esp, %ebp  
subl  $16, %esp  
movl  $0, -4(%ebp)  
movl  -4(%ebp), %eax  
addl  $7, %eax  
leave  
ret
```

X86: The De Facto Standard

- Extremely popular for desktop computers
- Alternatives
 - ARM: popular on mobile
 - MIPS: very simple
 - Itanium: ahead of its time
- CISC (complex instruction set computing)
 - Over 100 distinct opcodes in the set
- Register poor
 - Only 8 registers of 32-bits, only 6 are general-purpose
- Variable-length instructions
- Built of many backwards-compatible revisions
 - Many security problems preventable... in hindsight



Let's Dive in To X86!




Registers

← 32 bits →

General
purpose

EAX	AX	AH	AL
EBX	BX	BH	BL
ECX	CX	CH	CL
EDX	DX	DH	DL
ESI			
EDI			
ESP	(stack pointer)		
EBP	(base pointer)		

Process memory layout



`.text`

- Machine code of executable

`.data`

- Global initialized variables

`.bss`

- Below Stack Section
global uninitialized variables

heap

- Dynamic variables

stack

- Local variables
- Function call data

Env

- Environment variables
- Program arguments

Reminder: These are conventions

- Dictated by compiler
- Only instruction support by processor
 - Almost no structural notion of memory safety
 - Use of uninitialized memory
 - Use of freed memory
 - Memory leaks
- So how are they actually implemented?

Instruction Syntax

Examples:

```
subl $16, %ebx
```


```
movl (%eax), %ebx
```

opcode src, dst

- Constants preceded by **\$**
- Registers preceded by **%**
- Indirection uses **()**

Register Instructions: **sub**

registers




- Subtract from a register value

memory




```
subl %eax, %ebx
```

The Stack

- Local storage
 - Good place to keep data that doesn't fit into registers
- Grows from high addresses towards low addresses







Frame Instructions: **push**



- Put a value on the stack
 - Pull from register
 - Value goes to %esp
 - Subtract from %esp
- Example:


pushl %eax

Frame Instructions: pop







- Take a value from the stack
 - Pull from stack pointer
 - Value goes from %esp
 - Add to %esp

Control flow instructions: `jmp`



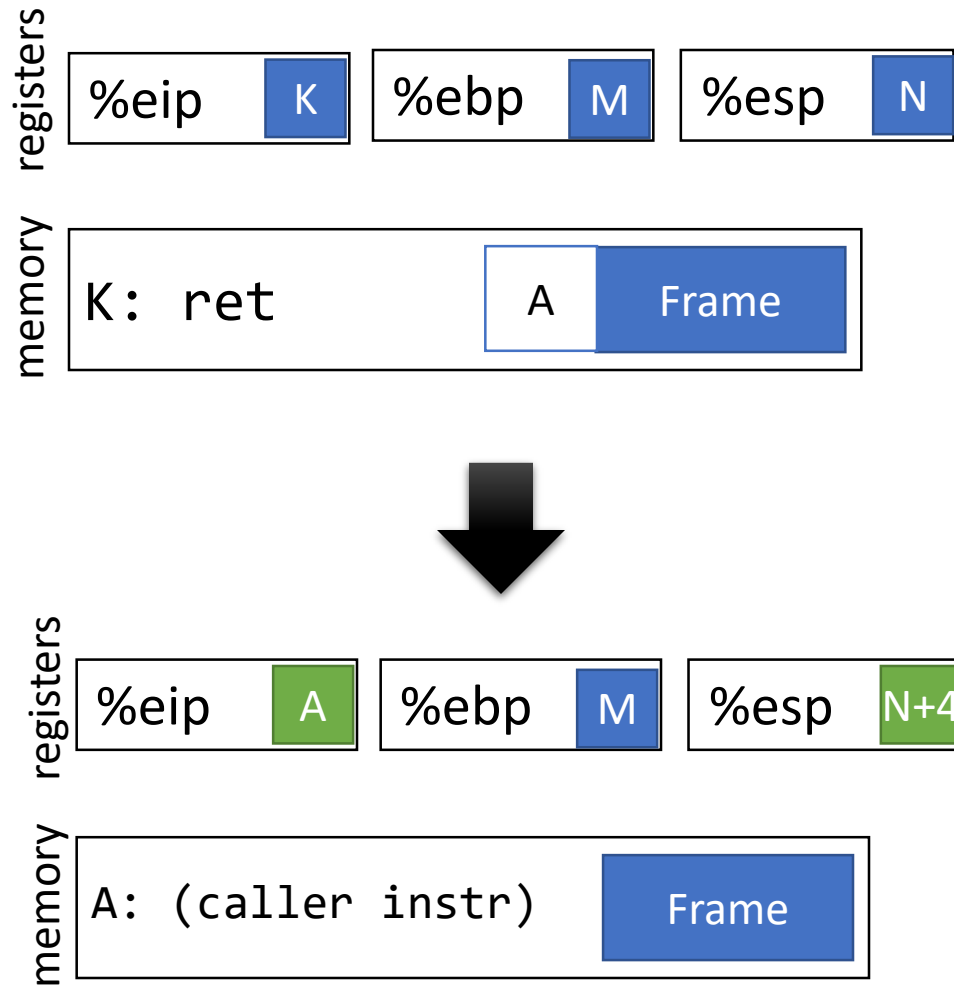
- `%eip` points to the currently executing instruction (in the text section)
- Has unconditional and conditional forms
- Uses relative addressing

Control flow instructions: `call`




- Saves the current instruction pointer to the stack
- Jumps to the argument value

Control flow instructions: `ret`



- Pops the stack into the instruction pointer

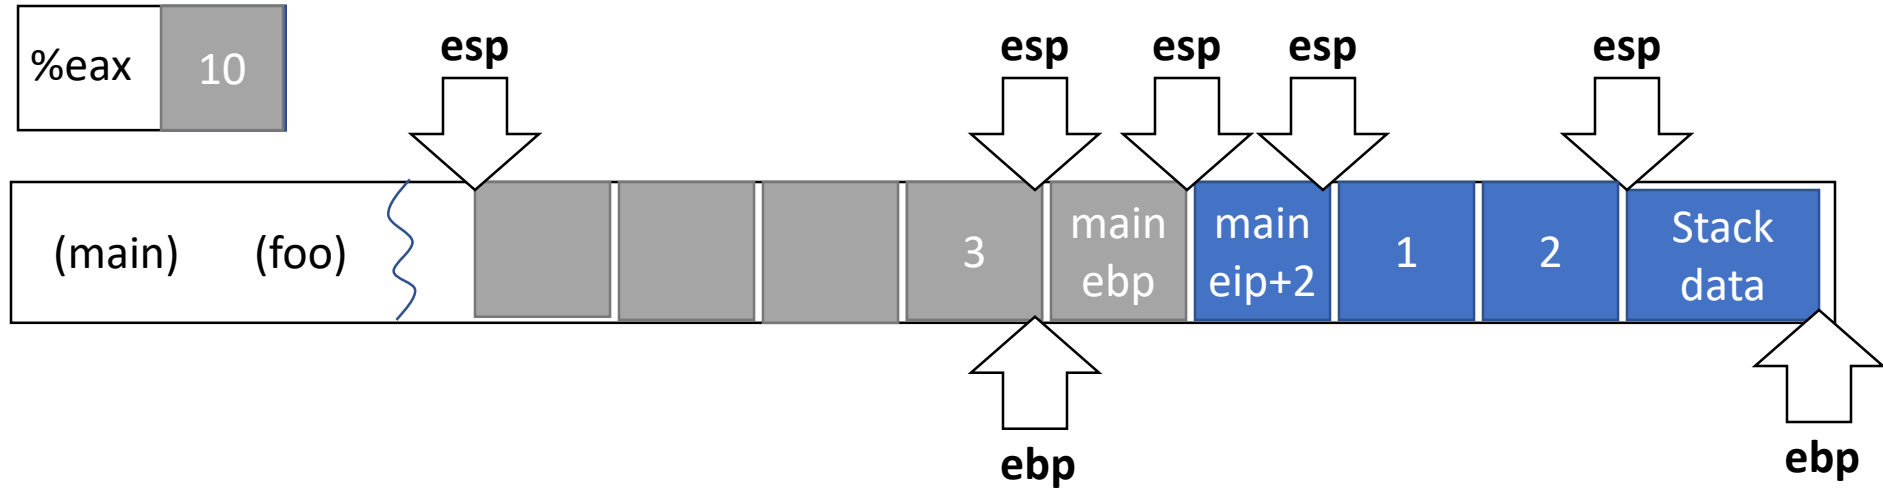
Stack instructions: **leave** (and **enter**)



- Equivalent to

```
movl %ebp, %esp  
popl %ebp
```
- copy EBP to ESP and then restore the old EBP from the stack

Implementing a function call



`main:`

```
...  
eip → subl    $8, %esp  
eip → movl    $2, 4(%esp)  
eip → movl    $1, (%esp)  
eip → call    foo  
eip → addl    $8, %esp  
...
```

`foo:`

```
eip → pushl   %ebp  
eip → movl    %esp, %ebp  
eip → subl    $16, %esp  
eip → movl    $3, -4(%ebp)  
eip → movl    8(%ebp), %eax  
eip → addl    $9, %eax  
eip → leave  
eip → ret
```


Function Calls: High level points

- Locals are organized into stack frames
 - Callees exist at lower address than the caller
- On call:
 - Save `%eip` so you can restore control
 - Save `%ebp` so you can restore data
- Implementation details are largely by convention
 - Somewhat codified by hardware

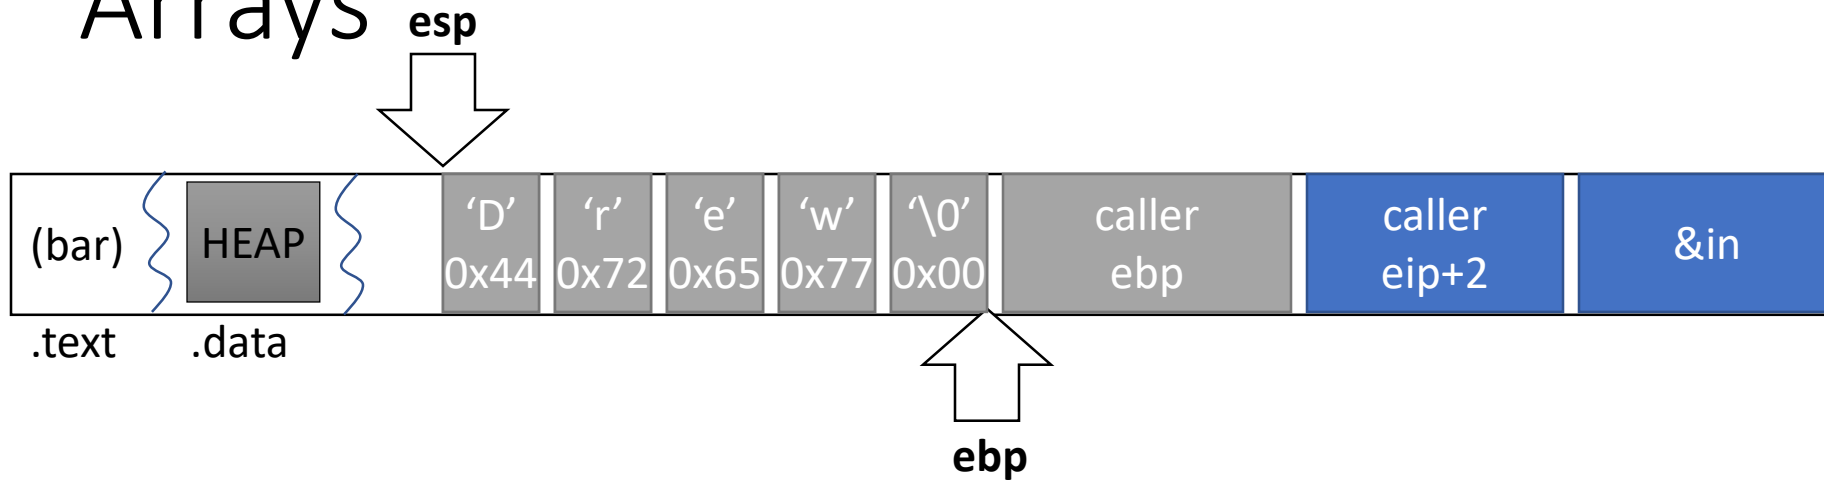
Data types / Endianness

- x86 is a little-endian architecture

`pushl %eax`



Arrays



```
void bar(char * in){  
    char name[5];  
    strcpy(name, in);  
}
```

```
bar:  
    pushl    %ebp  
    movl     %esp, %ebp  
    subl     $5, %esp  
    movl     8(%ebp), %eax  
    movl     %eax, 4(%esp)  
    leal     -5(%ebp), %eax  
    movl     %eax, (%esp)  
    call     strcpy  
    leave  
    ret
```

Tools: GCC

```
gcc -O0 -S program.c -o program.S -m32
```

Generate Assembly Code

```
gcc -O0 -g program.c -o program -m32
```

Generate Debugging Information

Tools: GDB

```
gdb program
```

```
(gdb) run
```

```
(gdb) list /* show the high-level code */
```

```
(gdb) decompile foo
```

```
(gdb) disas foo /* show assembly of foo */
```

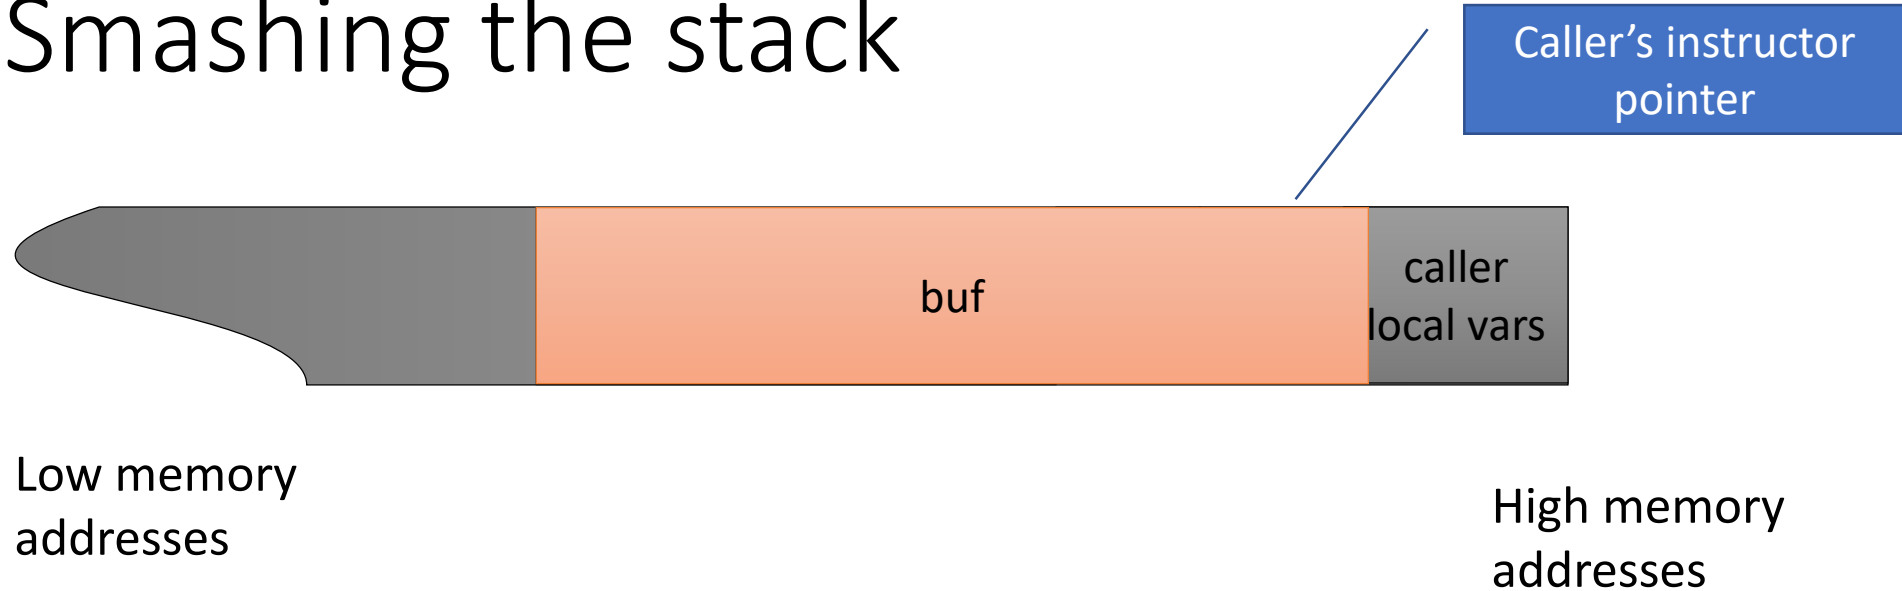
```
(gdb) disas main
```

```
(gdb) quit
```

x86 Summary

- Basics of x86
 - Process layout
 - ISA details
 - Most of the instructions that you'll need
- Introduced the concept of a buffer overflow
- Some tools to play around with x86 assembly

Smashing the stack



```
#include<stdio.h>
#include<string.h>

int main(int argc, char *argv[]) {
    char buf[100];
    strcpy(buf, argv[1]);
    printf("Hello %s\n", buf);
    return 0;
}
```

If argv[1] has more than 100 bytes...